

## Multiple choice questions

1. [0.5 point] Select all methods (of socket class from Python's socket module) that block the interpreter by default. For incorrect selection, some point will be deducted.
  - a. **accept()**
  - b. listen()
  - c. bind()
  - d. **recvfrom()**
2. [0.5 point] Choose all correct statements about physical clocks. For incorrect selection, some point will be deducted.
  - a. Time-of-day clock always moves forward
  - b. Monotonic clocks may move backwards
  - c. Better to use time-of-day clock to measure elapsed time on a single node
  - d. **Better to use monotonic clock to measure elapsed time on a single node**
3. [0.5 point] Choose all correct statements for logical clocks. For incorrect selection, some point will be deducted.
  - a. Given two Lamport timestamps such that  $L(a) < L(b)$ , then  $a \rightarrow b$  is always TRUE
  - b. **Given two Vector timestamps such that  $V(a) < V(b)$ , then  $a \rightarrow b$  is always TRUE**
  - c. **Lamport clocks implement partial order**
  - d. **Given two different vector timestamps we can surely say whether they are causal or concurrent**
4. [0.5 point] Choose all correct statements about multicast ordering. For incorrect selection, some point will be deducted.
  - a. **Causal ordering ensures FIFO ordering**
  - b. FIFO ordering ensures Causal ordering
  - c. Total ordering ensures Causal ordering
  - d. FIFO ordering ensures Total ordering

## Open questions

5. [1.5 points] Explain the FIFO, Causal, and Total ordered multicast; with up to three sentences for each of them.
  - a. **FIFO multicast**
    - If  $m1$  and  $m2$  are multicast messages by the same node, and  $m1$  is multicasted before  $m2$ , i.e.,  $\text{multicast}(m1) \rightarrow \text{multicast}(m2)$ , then  $m1$  must be delivered before  $m2$ , at all receivers
    - Multicasts by different nodes can be delivered at different order
  - b. **Causal multicast**
    - Multicast messages that are causally related must be delivered in the same causality order, at all receivers
      - If  $\text{multicast}(G, m1) \rightarrow \text{multicast}(G, m2)$
      - then every process that delivers  $m2$  will have delivered  $m1$
    - Concurrent multicast messages can be delivered in either order
  - c. **Total-ordered multicast**
    - All multicast messages are delivered at all group members in the same order
    - Unlike FIFO and Causal, this does not pay attention to order of multicast sending

**6. [2 points] Describe how following algorithms can be implemented.**

**a. [1 point] FIFO multicast. Explain the multicast sender and receiver operations.**

Key idea: put sequence number while sending messages

Each sender process

- embeds the sequence number (SN) into message
- increments the SN after sending the message
- may need a buffer to keep the copy of un-acknowledged messages

Each receiving process keeps track of SNs of all senders. When a message is received

- Reply with ACK message to inform about the reception
- If message SN is:
  - o as expected (next sequence), accept
  - o higher than expected, buffer in a queue
  - o lower than expected, reject

TCP implements FIFO ordering for unicast!

**b. [1 point] Total-ordered multicast using a Sequencer. Explain the multicast sender and receiver operations**

Sending multicast at process  $P_i$

- Instead of multicasting the message to everyone in group, the sender unicasts the message  $M$  to a special node called Sequencer

Sequencer


- A special node/process
- Maintains a global sequence number  $S$  (initially 0)
- When it receives a multicast message  $M$ , it sets  $S = S + 1$ , and multicasts  $\langle M, S \rangle$  to all processes in group

Receive multicast at process  $P_i$

- $P_i$  maintains a local received global sequence number  $S_i$  (initially 0)
- If  $P_i$  receives a multicast  $M$  from  $P_j$ , it buffers it until it both
  - o  $P_i$  receives  $\langle M, S(M) \rangle$  from sequencer, and
  - o  $S_i + 1 = S(M)$
- Then deliver it to application and set  $S_i = S_i + 1$


## 7. [2.5 points] Problems on logical clock

- a. [1 point] Put correct Lamport timestamps for following diagram



Event	Timestamp
a	1
b	2
c	4
d	5
e	2
f	3
g	1
h	6

- b. [1.5 points] Put correct Vector timestamps for following diagram



Event	Timestamp
a	1.0.0
b	2.0.0
c	3.2.0
d	4.2.0
e	1.1.0
f	1.2.0
g	0.0.1
h	4.2.2

## 8. [2 points] Problems on Chord protocol (2 points total)

- a. [1 point] Given the chord overlay with
- $m=5$
- and nodes [2, 16, 24, 26, 31], build the finger table for each of the nodes for
- $s=m$
- .

Node id	Node's finger table (list of the nodes in finger table)
2	16, 24
16	24, 2
24	26, 31, 2, 16
26	31, 2, 16
31	2, 16

- b. [1 point] Given the chord overlay and the finger tables (as above), find how the following lookup requests are resolved.

*Example:* Lookup(A) at node B? *Answer:* node B → node C → node D. Node D returns the address of node E (which is responsible for A)

i. Lookup(25) at node 2: 2 → 24, 24 returns addr of 25

ii. Lookup(0) at node 16: 16 → 24 → 31, 31 returns addr of 2